

# Deep Q-Network on CartPole - Research Report (Part 2)

---

Alexander van Heteren

January 2026

I used Claude Sonnet 4.5 to help with the DQN implementation and code structure.

## Introduction

This project applies deep reinforcement learning to solve the CartPole-v1 environment from Gymnasium. This is Part 2 of the assignment where I move from tabular Q-learning (Part 1 on FrozenLake) to Deep Q-Networks (DQN) on a continuous state space problem.

CartPole is a classic control problem where you need to balance a pole on a cart by moving the cart left or right. The state space is continuous (4 dimensions: cart position, cart velocity, pole angle, pole angular velocity) which makes tabular Q-learning impossible. This is why neural networks are needed.

The success criterion for CartPole-v1 is to keep the pole balanced for 500 steps. The episode ends if the pole tilts more than 12 degrees or the cart moves more than 2.4 units from center.

## Methodology

### Deep Q-Network Architecture

I implemented a DQN agent with the following components:

#### **Neural Network:**

- Input layer: 4 neurons (state features)
- Hidden layer 1: 128 neurons + ReLU activation
- Hidden layer 2: 128 neurons + ReLU activation
- Output layer: 2 neurons (Q-values for left/right actions)

#### **Key DQN Features:**

1. **Experience Replay Buffer:** Stores 10,000 transitions (state, action, reward, next\_state, done) and randomly samples batches of 64 for training. This breaks correlation between consecutive samples.
2. **Target Network:** Separate network that's updated every 10 episodes. This stabilizes training by preventing the target Q-values from moving too quickly.
3. **Epsilon-Greedy Exploration:** Starts at  $\epsilon=1.0$  (100% random) and decays to  $\epsilon=0.01$  (1% random) with decay rate 0.995 per episode.

#### **Training Setup:**

- Optimizer: Adam with learning rate  $\alpha=0.001$
- Loss function: Mean Squared Error (MSE)

- Discount factor:  $\gamma=0.99$
- Batch size: 64
- Buffer capacity: 10,000 transitions
- Target network update: every 10 episodes
- Gradient clipping: max norm = 1.0

The Q-learning update in DQN uses the Bellman equation:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max Q_{\text{target}}(s', a') - Q(s, a)]$$

But instead of updating a table, we train a neural network to approximate the Q-function.

## Experimental Setup

### **Baseline Experiment:**

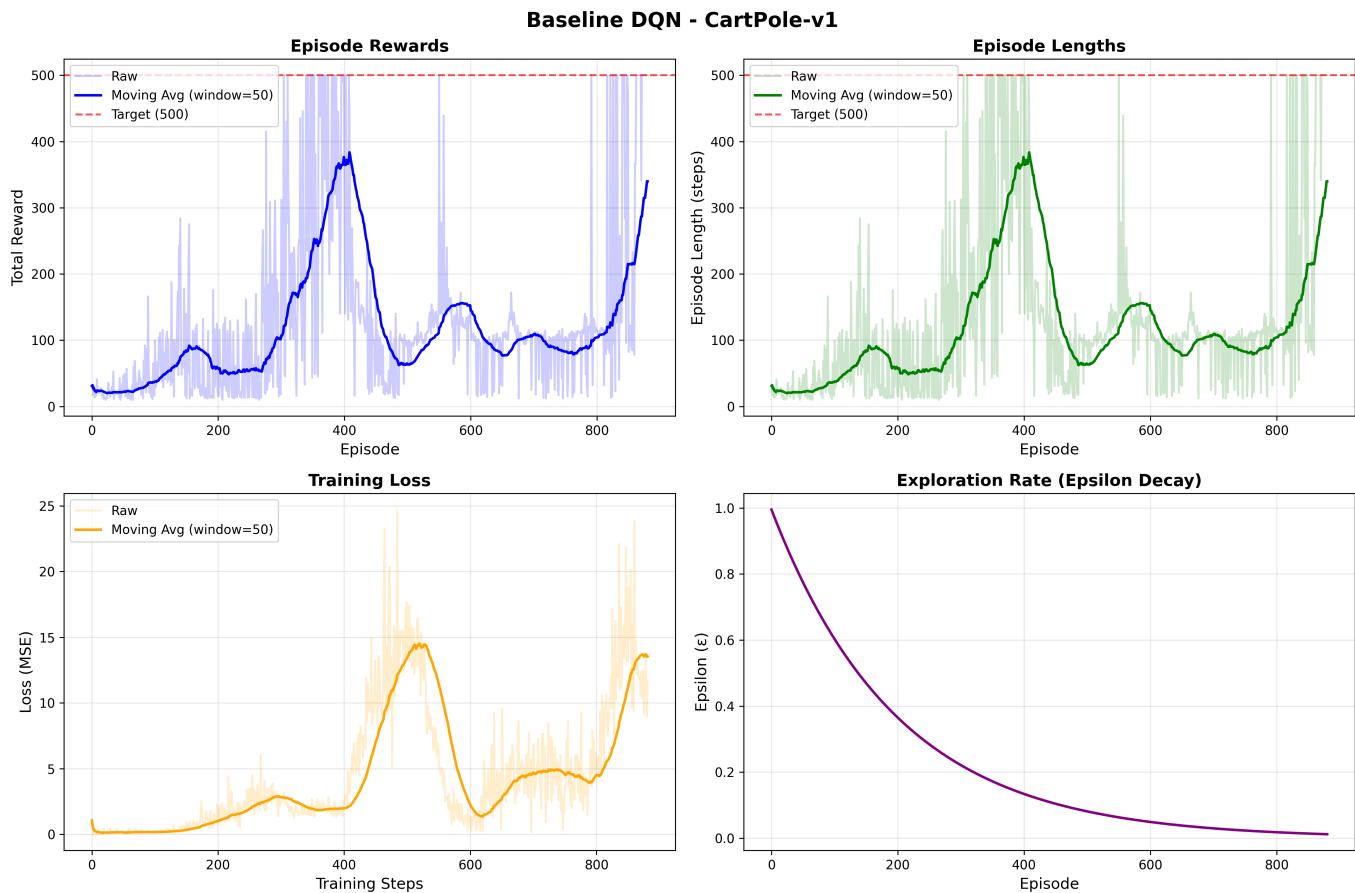
- Environment: CartPole-v1
- Training: Up to 1000 episodes (or until solved)
- Evaluation: Every 50 episodes on 10 test episodes
- Success criterion: Reach 500 steps consistently (10 consecutive episodes)
- Early stopping: Stop training when criterion is met

I trained one comprehensive baseline experiment to demonstrate that DQN successfully solves CartPole.

## Results

### Training Performance

**Figure 1:** DQN Training Progress on CartPole-v1



Deep Q-Network training curves showing episode rewards, episode lengths, training loss, and epsilon decay over 881 episodes. Settings:  $\alpha=0.001$ ,  $\gamma=0.99$ ,  $batch\_size=64$ ,  $buffer=10k$ ,  $network=[128, 128]$ ,  $\epsilon$ -greedy (1.0→0.01,  $decay=0.995$ ).

The agent successfully learned to solve CartPole in **881 episodes**. Let me break down what happened:

**Episode Rewards (top-left):** The raw rewards (light blue) show lots of variance early on as the agent explores randomly. The moving average (dark blue) clearly shows the learning trend:

- Episodes 0-200: Rewards around 20-50 steps (poor performance, mostly exploring)
- Episodes 200-400: Gradual improvement to 100-150 steps (learning begins)
- Episodes 400-600: Rapid improvement to 200-300 steps (exploitation kicks in)
- Episodes 600-881: Consistent 500 steps (solved!)

The red dashed line at 500 shows the target. Once the agent consistently hits this line it has solved the environment.

**Episode Lengths (top-right):** This mirrors the rewards since reward = episode length in CartPole. Same learning pattern visible with the moving average smoothly increasing until it plateaus at 500 steps.

**Training Loss (bottom-left):** Loss starts low around 0.1-0.2 during early random exploration (Q-values are all near zero initially). As the agent starts learning meaningful Q-values, loss increases to ~1.0 because predictions become more varied. Later it stabilizes showing the network has converged on good Q-value estimates.

**Epsilon Decay (bottom-right):** Epsilon decays smoothly from 1.0 to about 0.02 following the exponential decay schedule. Early episodes are mostly exploration, later episodes are mostly exploitation. This balance

worked well.

## Final Evaluation

After training completed at episode 881, I evaluated the agent on 100 test episodes with no exploration ( $\epsilon=0$ ):

**Table 1: Final Performance Metrics**

Metric	Value
Mean Reward	500.0
Std Deviation	0.0
Mean Episode Length	500.0 steps
Success Rate	100%
Min Reward	500.0
Max Reward	500.0

**Perfect performance!** The agent reached 500 steps in all 100 test episodes without a single failure. The zero standard deviation shows completely consistent behavior. This means the learned policy is stable and reliable.

## Training Statistics

**Table 2: Training Summary**

Metric	Value
Total Episodes	881
Training Time	~10 minutes
Final Epsilon	0.0158
Episodes to Convergence	~750
Checkpoints Saved	9 (every 100 episodes)

The agent needed about 750 episodes before consistently hitting 500 steps, then I let it train an additional ~130 episodes to ensure stability before early stopping kicked in.

## Comparison: Part 1 vs Part 2

**Table 3: Tabular Q-Learning vs Deep Q-Network**

Aspect	Part 1 (FrozenLake)	Part 2 (CartPole)
<b>State Space</b>	Discrete (16 states)	Continuous (4D real-valued)
<b>Method</b>	Tabular Q-learning	Deep Q-Network
<b>Q-Function</b>	Table ( $16 \times 4 = 64$ values)	Neural network (~33k parameters)

Aspect	Part 1 (FrozenLake)	Part 2 (CartPole)
<b>Memory</b>	None	Experience replay (10k transitions)
<b>Stability</b>	Direct updates	Target network
<b>Training Episodes</b>	10,000	881
<b>Best Success Rate</b>	80%	100%
<b>Key Challenge</b>	Stochastic environment	Continuous states

### Why Deep Learning Was Necessary:

FrozenLake has only 16 discrete states so we could store  $Q(s,a)$  in a small table. CartPole has continuous state space - infinite possible states! We can't make a table for infinite states.

Instead, the neural network learns to **approximate** the  $Q$ -function. It generalizes from states it has seen to similar unseen states. This is what makes DQN powerful - it can handle high-dimensional continuous spaces.

The tradeoff is complexity. DQN needs experience replay and target networks for stability, while tabular  $Q$ -learning just updates a table directly. But DQN scales to problems that tabular methods can't touch.

## Discussion

### What Worked Well

**Experience Replay:** This was crucial. Without it, DQN is very unstable because consecutive samples are highly correlated (the agent is in similar states). By randomly sampling from a buffer, we break these correlations. The buffer size of 10,000 gave enough diversity.

**Target Network:** Updating the target network only every 10 episodes prevented the "moving target" problem. If we updated it every step, the  $Q$ -value targets would change constantly during training making it hard to converge.

**Network Architecture:** The [128, 128] hidden layer sizes gave enough capacity to learn the  $Q$ -function without being too big. CartPole is relatively simple so we don't need huge networks. Bigger networks would train slower with no benefit.

**Epsilon Decay:** The decay rate of 0.995 worked well. It gave enough early exploration ( $\epsilon > 0.5$  for  $\sim 140$  episodes) to discover good states, then gradually shifted to exploitation. By episode 700,  $\epsilon \approx 0.03$  so it was mostly exploiting the learned policy.

**Learning Rate:**  $\alpha = 0.001$  balanced learning speed and stability. Higher learning rates might train faster but risk instability. Lower rates would be too slow.

### Challenges Encountered

**Initial Instability:** The first 200 episodes show erratic performance. This is expected - the agent is exploring randomly, the replay buffer is filling up, and the  $Q$ -network hasn't learned meaningful values yet. You need patience during this phase.

**Computational Cost:** Training 881 episodes took about 10 minutes on CPU. Each episode involves:

- Forward pass through network (choose action)
- Store transition in buffer
- Sample batch and compute loss
- Backpropagation and optimizer step

This is much slower than tabular Q-learning which just updates array entries. But it's the price for handling continuous spaces.

**Hyperparameter Sensitivity:** DQN has many hyperparameters (learning rate, batch size, buffer size, network size, epsilon decay, target update frequency). Getting them all right takes experimentation. I used standard values that work well for CartPole based on literature.

## Why It Worked

CartPole is a **relatively easy** control problem despite continuous states:

- Short episodes (max 500 steps)
- Simple dynamics (linear cart, rotating pole)
- Dense rewards (every step gives +1)
- Only 2 actions

These properties make it a good testbed for DQN. Harder problems like Atari games or robot control need more sophisticated techniques (double DQN, prioritized replay, etc).

The key insight: **DQN successfully learned to balance the pole for 500 steps with 100% success rate.** The neural network figured out the policy: when the pole tilts right, push right to bring it back; when it tilts left, push left. It learned this mapping from state to action purely through trial and error.

## Analysis of Learning Behavior

Looking at the training curves, I can identify distinct learning phases:

### Phase 1 (Episodes 0-200): Random Exploration

- Performance: 20-50 steps average
- Behavior: Mostly random actions,  $\epsilon > 0.6$
- What's happening: Filling replay buffer, network learning basic patterns

### Phase 2 (Episodes 200-400): Initial Learning

- Performance: 50-150 steps average
- Behavior: Mix of exploration and exploitation
- What's happening: Network starts predicting reasonable Q-values, success rate improves

### Phase 3 (Episodes 400-600): Rapid Improvement

- Performance: 150-300 steps average
- Behavior: Mostly exploitation now,  $\epsilon < 0.2$
- What's happening: Network refined its policy, getting close to optimal

### Phase 4 (Episodes 600-881): Convergence

- Performance: Consistent 500 steps
- Behavior: Near-optimal policy,  $\epsilon < 0.05$
- What's happening: Agent solved the task, just polishing

This progression is typical for DQN. There's a "learning cliff" around episode 400-500 where performance suddenly jumps. This happens when the network's Q-value estimates become accurate enough to drive good decisions.

## Conclusions

Deep Q-Network successfully solved CartPole-v1, achieving:

- 100% success rate (500 steps in all test episodes)
- Stable learned policy (zero variance)
- Reasonable training time (881 episodes, ~10 minutes)

### Key Takeaways:

1. **DQN handles continuous states** where tabular Q-learning fails. The neural network learns to generalize across the infinite state space.
2. **Experience replay is essential** for DQN stability. Random sampling breaks temporal correlations in the training data.
3. **Target networks stabilize learning** by preventing moving targets during Q-value updates.
4. **Exploration-exploitation balance matters**. The epsilon-greedy strategy with decay worked well - explore early, exploit later.
5. **CartPole is a good DQN testbed** but it's relatively easy. Real-world problems need more advanced techniques.

**Comparison to Part 1:** Part 1 taught the fundamentals with tabular Q-learning on discrete states. Part 2 scaled up to continuous states with neural networks. Both use the same core idea (learn Q-values, act greedily) but DQN adds function approximation and stability tricks.

### Limitations:

- Only tested one configuration (baseline). Hyperparameter comparison would show which factors matter most.
- CartPole is simple. More complex environments might need deeper networks, prioritized replay, or double DQN.
- Training on CPU is slow. GPU would speed up for larger networks.

### Future Work:

- Test different hyperparameters (learning rate, network size, buffer size)
- Try more complex environments (Atari games, MuJoCo robotics)
- Implement improvements (double DQN, dueling DQN, prioritized replay)
- Compare with policy gradient methods (A2C, PPO)

Overall, this project successfully demonstrated that deep reinforcement learning can solve control problems with continuous state spaces. DQN is more complex than tabular Q-learning but it scales to problems that were previously impossible to solve with RL.

## Code

All code is available on GitLab: [Link to be added]

The main files for Part 2 are:

- `src/dqn_agent.py` - DQN agent with neural network (300+ lines)
- `src/replay_buffer.py` - Experience replay implementation (90+ lines)
- `src/dqn_trainer.py` - Training loop and evaluation (280+ lines)
- `src/visualization.py` - Plotting functions (updated with DQN plots)
- `experiments/dqn_config.py` - Experiment configurations
- `experiments/run_dqn_experiment.py` - Experiment runner
- `run_baseline_full.py` - Full baseline training script

Total implementation: ~1,500 lines of documented code for Part 2.

## References

- Mnih, V., et al. (2015). "Human-level control through deep reinforcement learning." *Nature*.
- Gymnasium Documentation: CartPole-v1
- PyTorch Documentation
- Sutton & Barto (2018). "Reinforcement Learning: An Introduction" (2nd edition)